

Introduction to the New Statistics - R Workbook

Authors here

16 May 2016

Contents

1	Overview	3
2	R Basics	4
2.1	Useful Websites for Learning R	4
3	Picturing and Describing Data	5
3.1	Descriptive Statistics	5
3.2	Examples	5
3.3	Summary function	6
3.4	Statitics by group	7
3.5	Apply to multiple variables at a time	8
3.6	Histograms	9
3.7	Stacked Dot Plots	12
3.8	Multiple groups	16
3.9	Clicable links to other sources of information	17
3.10	Other Links	17
4	Z-score and T-scores	17
4.1	Compute z-scores	17
4.2	Find a tail probability from a raw score	19
4.3	Find a tail probabilities from z-scores	19
4.4	Find Z-scores given tail probabilities	20
4.5	Find raw scores given tail probabilities	20
4.6	Find tail areas for chosen t	20
4.7	Find critical t-values given df and tail area	20
4.8	Calculate p-values given z-scores or t-scores	21
5	Effect sizes and Confidence Intervals for Means and Mean Differences	21
5.1	Confidence Interval for the Mean of a Single Sample	21
5.2	Two Independent Groups	22
5.3	Trimmed mean difference	24
5.4	Two Dependent Groups	24
5.5	One and two-way designs	24

6	Correlation	25
6.1	Figures	25
7	Regression	25
8	Frequencies, Proportions, Risks	27
9	Precision and Planning	27
10	To do	27
10.1	Cat's Eye Picture	27
10.2	MOE for means ?	27

1 Overview

An overview of the workbook to go here.

Info on how to install R functions etc.

Load the `itns` library. This will make all the data sets used in the book accessible.

```
library(itns)
```

2 R Basics

Some basic info for R beginners to go here.

eg how to install R, and sources of help.

This section can be skipped by anyone who already knows how to use R.

2.1 Useful Websites for Learning R

A list of useful websites, online tutorials etc for learning R to go here.

R-Studio Resources for Learning R (with links to interactive tutorials for beginners)

R Bloggers - How to Learn R

Quick-R

Cookbook for R

R Quick Reference Card

List a few particularly useful packages

dplyr

https://www3.nd.edu/~steve/computing_with_data/24_dplyr/dplyr.html

rehape and reshape2

etc

3 Picturing and Describing Data

Materials relevant to ITNS Ch 3 here.

3.1 Descriptive Statistics

List of inbuilt R functions to compute basic summary statistics.

Measure	R function
Mean	<code>mean()</code>
Median	<code>median()</code>
Minimum	<code>min()</code>
Maximum	<code>max()</code>
Range	<code>range()</code>
Interquartile Range	<code>IQR()</code>
Variance	<code>var()</code>
Standard Deviation	<code>sd()</code>
Percentiles	<code>quantile()</code>

3.2 Examples

Using the pen laptop data.

Here are the first few cases:

```
head(pen_laptop1)
```

```
##   group transcription
## 1   Pen           12.1
## 2   Pen             6.5
## 3   Pen             8.1
## 4   Pen             7.6
## 5   Pen           12.2
## 6   Pen           10.8
```

Compute summary statistics:

```
# Minimum
min(pen_laptop1$transcription)
```

```
## [1] 1
```

```
# Maximum
max(pen_laptop1$transcription)
```

```
## [1] 34.7
```

```
# Mean  
mean(pen_laptop1$transcription)
```

```
## [1] 11.53385
```

```
# Median  
median(pen_laptop1$transcription)
```

```
## [1] 10.7
```

```
# 95th percentile  
quantile(pen_laptop1$transcription, probs = .95)
```

```
## 95%  
## 21.34
```

```
# From the 5th to 95th percentiles, in steps of 5  
quantile(pen_laptop1$transcription, probs = seq(from = .05, to = .95, by = .05))
```

```
## 5% 10% 15% 20% 25% 30% 35% 40% 45% 50% 55% 60%  
## 2.16 3.38 5.10 6.24 8.00 8.50 8.70 9.16 9.68 10.70 11.22 12.04  
## 65% 70% 75% 80% 85% 90% 95%  
## 12.64 13.20 15.20 17.06 17.82 18.92 21.34
```

```
# Variance  
var(pen_laptop1$transcription)
```

```
## [1] 44.7654
```

```
# Standard deviation  
sd(pen_laptop1$transcription)
```

```
## [1] 6.690695
```

```
# Range  
range(pen_laptop1$transcription)
```

```
## [1] 1.0 34.7
```

```
# Interquartile range  
IQR(pen_laptop1$transcription)
```

```
## [1] 7.2
```

3.3 Summary function

```
summary(pen_laptop1$transcription)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00   8.00   10.70   11.53   15.20   34.70
```

3.4 Statistics by group

Using inbuilt by() function

```
by(pen_laptop1$transcription, pen_laptop1$group, summary)
```

```
## pen_laptop1$group: Laptop
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.20   9.45   12.80   14.52   17.85   34.70
## -----
## pen_laptop1$group: Pen
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.000   5.200   8.600   8.812  11.280  20.100
```

Using the more powerful and flexible dplyr package

```
# Load the dplyr package
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##      filter, lag

## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
```

```
group_by(pen_laptop1, group) %>%
  summarize(avg = mean(transcription))
```

```
## Source: local data frame [2 x 2]
##
##      group      avg
##   (fctr)   (dbl)
## 1 Laptop 14.519355
## 2   Pen  8.811765
```

```
# Going further - find the mean and standard deviation for each group
group_by(pen_laptop1, group) %>%
  summarize(avg = mean(transcription),
            sd = sd(transcription)
            )
```

```
## Source: local data frame [2 x 3]
##
##   group      avg      sd
##   (fctr)    (dbl)    (dbl)
## 1 Laptop 14.519355 7.285576
## 2   Pen  8.811765 4.749339
```

For a two-way design, blame1 dataset

There are two independent variables - *Socioeconomic status (ses)* (high or low) and *race* (black or white).

```
blame1 %>%
  group_by(race, ses) %>%
  summarize(avg = mean(blame), std.dev = sd(blame))
```

```
## Source: local data frame [4 x 4]
## Groups: race [?]
##
##   race    ses      avg std.dev
##   (fctr) (fctr)  (dbl)  (dbl)
## 1 black  high 2.570423 1.842544
## 2 black  low 2.966292 1.824389
## 3 white  low 2.800000 1.734270
## 4 White  high 3.181818 1.837054
```

3.5 Apply to multiple variables at a time

summary() function does this by default.

Example: college_survey1 data

```
summary(college_survey1)
```

```
##      id          gender      age      school_year
## Min.   :21363617  Female:175  Min.    :18.00  First-year:43
## 1st Qu.:21442046  Male  : 68  1st Qu.:19.00  Sophomore :64
## Median :21466849                Median :20.00  Junior    :53
## Mean   :21456592                Mean   :21.79  Senior    :57
## 3rd Qu.:21470366                3rd Qu.:22.00  Post-bac  :26
## Max.   :21516638                Max.   :59.00
##                                     NA's    :4
## transfer      student_athlete student_athlete_code wealth_sr
## No  :209  In season   : 12  No  :212                Min.    :1.0
## Yes : 33  non-athlete :212  Yes : 30                1st Qu.:2.0
## NA's: 1  Out of season: 18  NA's: 1                Median :3.0
##                                     Mean    :2.9
##                                     3rd Qu.:3.0
##                                     Max.    :5.0
##                                     NA's    :3
##      gpa      act      subjective_well_being positive_affect
## Min.   :0.800  Min.    :15.00  Min.    :1.000  Min.    :1.200
## 1st Qu.:3.000  1st Qu.:21.00  1st Qu.:4.125  1st Qu.:3.000
## Median :3.400  Median :24.00  Median :5.000  Median :3.400
```



```
## Mean :3.343 Mean :24.18 Mean :4.941 Mean :3.433
## 3rd Qu.:3.800 3rd Qu.:27.00 3rd Qu.:6.000 3rd Qu.:3.900
## Max. :4.000 Max. :36.00 Max. :7.000 Max. :5.000
## NA's :6 NA's :28 NA's :8
## negative_affect relationship_confidence exercise
## Min. :1.000 Min. :1.000 Min. : 0.00
## 1st Qu.:1.700 1st Qu.:3.105 1st Qu.: 2.75
## Median :2.200 Median :3.770 Median : 22.00
## Mean :2.322 Mean :3.664 Mean : 54.78
## 3rd Qu.:2.800 3rd Qu.:4.290 3rd Qu.: 61.00
## Max. :4.600 Max. :5.000 Max. :1810.00
## NA's :8 NA's :11 NA's :23
## academic_motivation_intrinsic academic_motivation_extrinsic
## Min. :1.500 Min. :2.500
## 1st Qu.:4.000 1st Qu.:5.170
## Median :4.830 Median :6.000
## Mean :4.862 Mean :5.764
## 3rd Qu.:5.830 3rd Qu.:6.500
## Max. :7.000 Max. :7.000
## NA's :20 NA's :20
## academic_motivation_amotivation intelligence_value raven_score
## Min. :1.000 Min. :2.330 Min. : 0.00
## 1st Qu.:1.000 1st Qu.:3.330 1st Qu.:25.00
## Median :1.000 Median :3.670 Median :37.50
## Mean :1.971 Mean :3.619 Mean :37.71
## 3rd Qu.:2.500 3rd Qu.:4.000 3rd Qu.:50.00
## Max. :7.000 Max. :5.000 Max. :87.50
## NA's :20 NA's :20 NA's :26
```

Using dplyr package.

Split by gender, then find mean for three variables (subjective well being, positive affect, negative affect)

```
college_survey1 %>%
  group_by(gender) %>%
  summarise_each(funs(mean(., na.rm = TRUE))), subjective_well_being, positive_affect, negative_affect)
```

```
## Source: local data frame [2 x 4]
##
##   gender subjective_well_being positive_affect negative_affect
##   (fctr)           (dbl)           (dbl)           (dbl)
## 1 Female           4.903314           3.346           2.374176
## 2 Male             5.038088           3.662           2.185692
```


See also plyr, apply functions.

useful Links Quick R Guide to Basic Statistics in R

3.6 Histograms


3.6.1 Single Sample

```
library(lattice)
histogram(~transcription, data = pen_laptop1)
```




Alter the number of bins

```
histogram(~transcription, data = pen_laptop1, nint = 10)
```



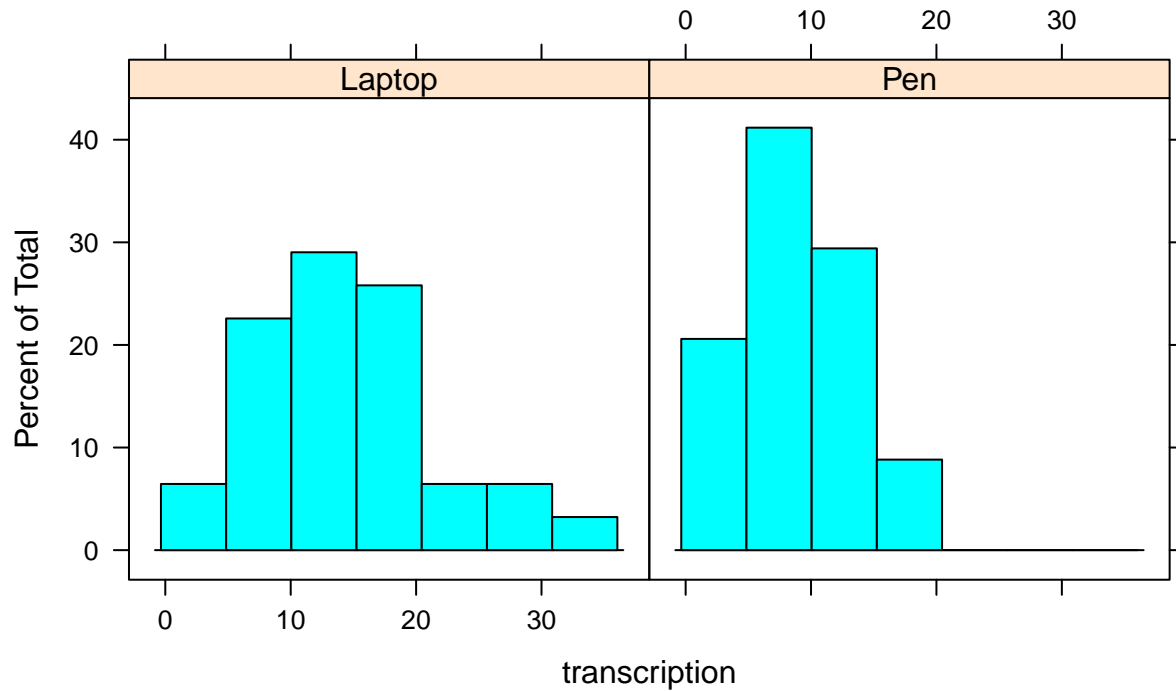
Display count rather than count on the Y-axis

```
histogram(~transcription, data = pen_laptop1, nint = 10, type = "count")
```



3.6.2 Multiple Samples

```
histogram(~transcription | group, data = pen_laptop1)
```



3.6.3 Links

Lattice Histogram Vignette


3.7 Stacked Dot Plots

3.7.1 Single group

A basic stacked dotplot generated using the `ggplot2` package.


```
library(ggplot2)
p <- ggplot(pen_laptop1, aes(x = transcription)) +
  geom_dotplot()
p
```

stat_bindot: binwidth defaulted to range/30. Use 'binwidth = x' to adjust this.




Change the binwidth, dotsize, and stack ratio

```
p <- ggplot(pen_laptop1, aes(x = transcription)) +  
  geom_dotplot(binwidth = 1.5, dotsize = 1.2, stackratio = 1.5)  
p
```




Center the dots

```
p <- ggplot(pen_laptop1, aes(x = transcription)) +
  geom_dotplot(binwidth = 1, dotsize = 1.2, stackratio = 1.5, stackdir = "center")
p
```




Rotate

```
p + coord_flip()
```



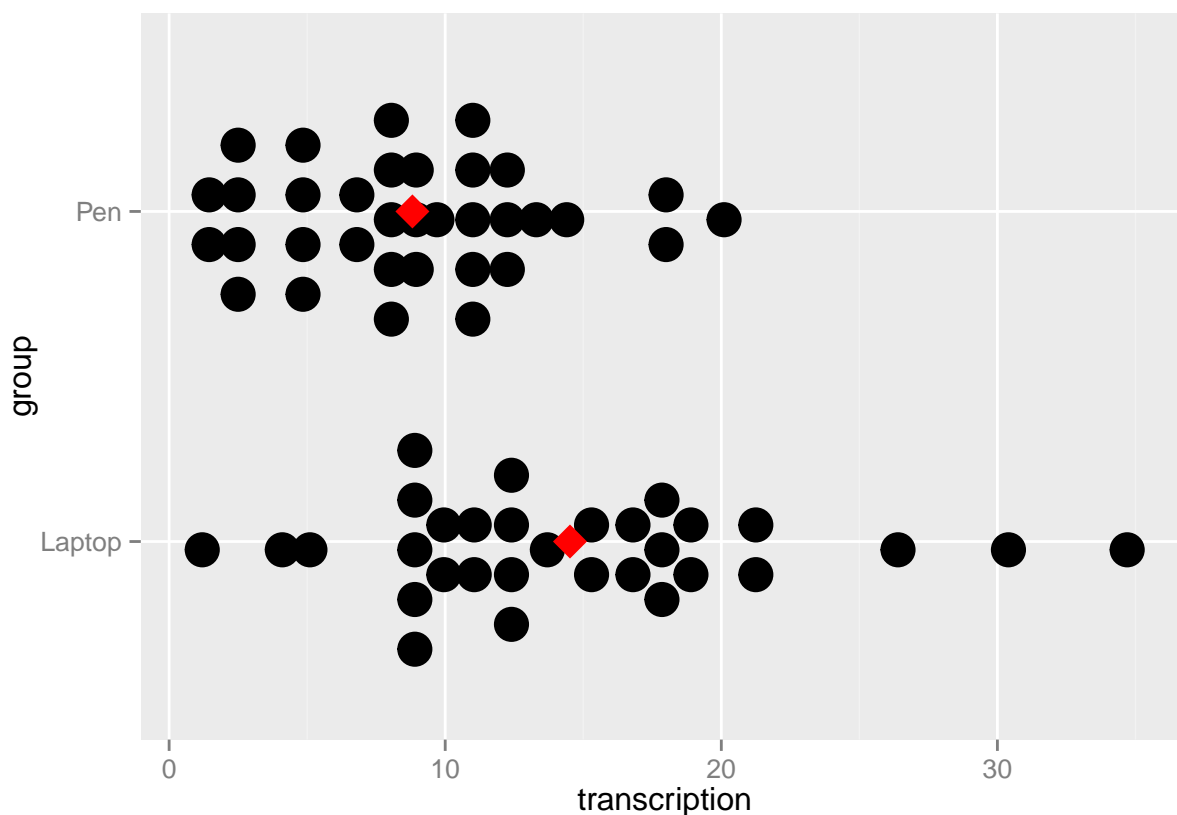
3.8 Multiple groups

```
p <- ggplot(pen_laptop1, aes(x = group, y = transcription)) +  
  geom_dotplot(binaxis = "y", binwidth = 1, dotsize = 1.2, stackratio = 1.5, stackdir = "center")  
p + coord_flip()
```



Add the mean of each group (as a red diamond)

```
p + stat_summary(fun.y = mean, geom = "point", shape = 18,  
  size = 6, color = "red") +  
  coord_flip()
```

3.9 Clickable links to other sources of information

- [ggplot2 dotplot tutorial on the Statistical tools for high-throughput data analysis website.](#)
- [Examples of how to generate dotplots using other R packages.](#)

3.10 Other Links

- [Examples of Lattice Graphics](#)

To add: Add lattice book and vignettes

ggplot2 [ggplot2 book](#) [ggplot2 website](#) [r bloggers](#) [ggplot2 cookbook](#)

4 Z-score and T-scores

4.1 Compute z-scores

4.1.1 `scale()`

The inbuilt `scale()` function converts raw scores to Z-scores. The sample mean is subtracted from each score, and the resulting values are divided by the sample standard deviation.

```
# For the pen_laptop dataset, convert all raw transcription scores to Z-scores  
scale(pen_laptop1$transcription)
```

```
##           [,1]  
## [1,]  0.08461809  
## [2,] -0.75236522  
## [3,] -0.51322713  
## [4,] -0.58795778  
## [5,]  0.09956422  
## [6,] -0.10968160  
## [7,] -1.57440239  
## [8,] -1.29042591  
## [9,]  0.42837909  
## [10,] -0.46838874  
## [11,]  0.92160140  
## [12,]  1.28030853  
## [13,] -1.40999496  
## [14,] -0.06484321  
## [15,] -0.04989708  
## [16,] -0.12462773  
## [17,] -1.43988722  
## [18,] -0.94666491  
## [19,] -0.27408904  
## [20,] -0.94666491  
## [21,] -1.36515657  
## [22,] -0.66268843  
## [23,] -0.42355034  
## [24,] -0.52817326  
## [25,] -0.03495095  
## [26,] -0.45344261  
## [27,] -0.36376582  
## [28,] -1.05128783  
## [29,] -0.34881969  
## [30,]  0.26397166  
## [31,]  1.01127818  
## [32,] -1.30537204  
## [33,] -0.96161104  
## [34,]  0.12945648  
## [35,]  0.32375618  
## [36,]  1.42976984  
## [37,]  0.54794814  
## [38,]  2.81975997  
## [39,]  0.18924101  
## [40,] -0.28903517  
## [41,] -0.33387356  
## [42,]  0.92160140  
## [43,]  0.57784040  
## [44,] -0.42355034  
## [45,]  0.18924101  
## [46,] -0.13957386  
## [47,] -0.96161104  
## [48,]  0.77214010  
## [49,]  0.92160140
```

```
## [50,] -0.42355034
## [51,]  2.22191475
## [52,]  0.96643979
## [53,]  1.11590110
## [54,]  0.80203236
## [55,]  1.08600884
## [56,] -0.45344261
## [57,] -1.54451013
## [58,] -0.00505869
## [59,]  1.47460823
## [60,] -0.18441226
## [61,] -0.37871195
## [62,]  0.18924101
## [63,]  0.06967196
## [64,]  3.46244359
## [65,] -1.11107235
## attr("scaled:center")
## [1] 11.53385
## attr("scaled:scale")
## [1] 6.690695
```

4.2 Find a tail probability from a raw score

Use IQ example from Ch 4. IQ tests have a population mean of 100 and SD of 15. Therefore a person with a score of 115 would have a z score of 1, and score higher than about 84% of the population.

Use the inbuilt `pnorm()` function to find that probability.

```
# Expressed on a 0 - 1 scale
pnorm(q = 115, mean = 100, sd = 15)
```

```
## [1] 0.8413447
```

```
# Expressed on 0 - 100 scale
pnorm(q = 115, mean = 100, sd = 15)* 100
```

```
## [1] 84.13447
```

4.3 Find a tail probabilities from z-scores

```
# Z-score of 1
pnorm(q = 1)* 100
```

```
## [1] 84.13447
```

```
# Z-scores of -1.96 and 1.96
pnorm(q = c(-1.96, 1.96))* 100
```

```
## [1]  2.49979 97.50021
```

There is no need to specify the mean and standard deviation as the defaults are 0 and 1.

4.4 Find Z-scores given tail probabilities

```
# Find Z-scores corresponding to area under the curve
```

```
qnorm(p = .025) # lower tail
```

```
## [1] -1.959964
```

```
qnorm(p = .025, lower.tail = FALSE) # upper tail
```

```
## [1] 1.959964
```

4.5 Find raw scores given tail probabilities

```
# Find raw scores corresponding to area under the curve
```

```
qnorm(p = .025, mean = 100, sd = 15) # lower tail
```

```
## [1] 70.60054
```

```
qnorm(p = .025, mean = 100, sd = 15, lower.tail = FALSE) # upper tail
```

```
## [1] 129.3995
```

4.6 Find tail areas for chosen t

e.g., t-value of 2.145 and $df = 14$

```
pt(q = 2.145, df = 14) # upper tail
```

```
## [1] 0.9750099
```

```
pt(q = -2.145, df = 14) # lower tail
```

```
## [1] 0.02499008
```

```
pt(q = 2.145, df = 14, lower.tail = FALSE) # another to find lower tail probability
```

```
## [1] 0.02499008
```

4.7 Find critical t-values given df and tail area

```
qt(p = .025, df = 14) # lower tail critical value
```

```
## [1] -2.144787
```

```
qt(p = .025, df = 14, lower.tail = FALSE) # upper tail critical value
```

```
## [1] 2.144787
```

```
qt(p = .975, df = 14) # alternate way to find upper tail critical t value
```

```
## [1] 2.144787
```

4.8 Calculate p-values given z-scores or t-scores

Tutorial here

5 Effect sizes and Confidence Intervals for Means and Mean Differences

5.1 Confidence Interval for the Mean of a Single Sample

Using t-distribution. There is no inbuilt function for Z as far as I know.

By default a 95% confidence interval for the mean is returned, as well as the estimated mean, t-value, degrees of freedom and p-value.

```
t.test(pen_laptop1$transcription)
```

```
##  
## One Sample t-test  
##  
## data: pen_laptop1$transcription  
## t = 13.898, df = 64, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 0  
## 95 percent confidence interval:  
## 9.875973 13.191719  
## sample estimates:  
## mean of x  
## 11.53385
```

To use a different confidence level, use the `conf.level` argument. For example, to compute a 99% confidence interval you would use

```
t.test(pen_laptop1$transcription, conf.level = .99)
```

```
##
## One Sample t-test
##
## data: pen_laptop1$transcription
## t = 13.898, df = 64, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 99 percent confidence interval:
## 9.330639 13.737053
## sample estimates:
## mean of x
## 11.53385
```

5.2 Two Independent Groups

5.2.1 Mean Difference and Confidence Interval

5.2.1.1 Assuming unequal variances

This is the default in R.

```
t.test(transcription ~ group, data = pen_laptop1)
```

```
##
## Welch Two Sample t-test
##
## data: transcription by group
## t = 3.7031, df = 50.816, p-value = 0.0005254
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 2.612991 8.802189
## sample estimates:
## mean in group Laptop    mean in group Pen
##      14.519355          8.811765
```

5.2.1.2 Assuming equal variances

```
t.test(transcription ~ group, data = pen_laptop1, var.equal = TRUE)
```

```
##
## Two Sample t-test
##
## data: transcription by group
## t = 3.7738, df = 63, p-value = 0.0003579
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 2.685265 8.729915
## sample estimates:
## mean in group Laptop    mean in group Pen
##      14.519355          8.811765
```

5.2.2 Standardized Mean Difference

Use the MBESS package.

Load the package, and then extract the data to be used in the analysis.

```
library(dplyr)
library(MBESS)

# Extract transcription scores for the 'Laptop' group
x <- pen_laptop1 %>% filter (group == "Laptop")

# Extract transcription scores for the 'Pen' group
y <- pen_laptop1 %>% filter (group == "Pen")

# Find the sample size for each group, excluding any missing data
nx <- na.omit(length(x$transcription))
ny <- na.omit(length(y$transcription))

## Find the standardized mean difference (biased)
d <- smd(x$transcription, y$transcription)
d

## [1] 0.9371681

## Find unbiased standardized mean difference
du <- smd(x$transcription, y$transcription, Unbiased = TRUE)
du

## [1] 0.9259595

# Compute a confidence interval
ci.smd (smd = d, n.1 = nx, n.2 = ny)

## $Lower.Conf.Limit.smd
## [1] 0.4204238
##
## $smd
## [1] 0.9371681
##
## $Upper.Conf.Limit.smd
## [1] 1.447208
```

The `smd` and `'ci.smd'` functions used the pooled standard deviation to compute `d`.

To use the standard deviation of one of the groups instead (e.g., when the homogeneity of variance assumption is not met), use the alternative functions `smd.c` and `'ci.smd.c'`

For example, to use the laptop group standard deviation as the standardizer:

```
# Biased d, using the Laptop SD as the standardizer
d <- smd.c(Group.T = y$transcription, # Pen raw data
           Group.C = x$transcription # Laptop raw data
           )
d
```

```
## [1] -0.7834096

# Confidence Interval
ci.smd.c (smd.c = d, n.C = nx, n.E = ny)
```

```
## $Lower.Conf.Limit.smd.c
## [1] -1.303023
##
## $smd.c
## [1] -0.7834096
##
## $Upper.Conf.Limit.smd.c
## [1] -0.2525728
```

5.3 Trimmed mean difference

Use WRS2 package.

5.4 Two Dependent Groups

5.4.1 Unstandardized Mean Change

Use thomason1 dataset.

```
t.test(thomason1$pre, thomason1$post, paired = TRUE)

##
## Paired t-test
##
## data: thomason1$pre and thomason1$post
## t = -3.8555, df = 11, p-value = 0.002674
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -2.618115 -0.715218
## sample estimates:
## mean of the differences
## -1.666667
```

5.4.2 Standardized Mean Change

There is no MBESS function for this.

Need to write a function or show calculations.

5.5 One and two-way designs

One-way independent groups

one way dependent groups

Two-way independent groups

Between by within

Contrasts

6 Correlation

6.1 Figures

scatterplot and other plots in R

corplot package

6.1.1 Correlations and Confidence interval

```
cor(thomason1)
```

```
##           pre      post
## pre  1.0000000 0.8923908
## post 0.8923908 1.0000000
```

```
cor.test(thomason1$pre, thomason1$post)
```

```
##
## Pearson's product-moment correlation
##
## data:  thomason1$pre and thomason1$post
## t = 6.2535, df = 10, p-value = 9.462e-05
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.6528351 0.9696774
## sample estimates:
##           cor
## 0.8923908
```

6.1.2 Difference in correlations

here. Have to look up.

7 Regression

Using home_prices data.

```
mod <- lm(price ~ size, data = home_prices)
mod
```

```
##
## Call:
## lm(formula = price ~ size, data = home_prices)
##
## Coefficients:
## (Intercept)      size
##    -66024      2713
```

```
summary(mod)
```

```
##
## Call:
## lm(formula = price ~ size, data = home_prices)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -627083  -90490  -34039   49459   955266
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -66024.3     23426.5  -2.818  0.00515 **
## size         2713.4         125.4   21.646 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 179100 on 298 degrees of freedom
## Multiple R-squared:  0.6112, Adjusted R-squared:  0.6099
## F-statistic: 468.6 on 1 and 298 DF,  p-value: < 2.2e-16
```

```
confint(mod)
```

```
##              2.5 %      97.5 %
## (Intercept) -112126.598 -19921.980
## size         2466.688    2960.059
```

```
anova(mod)
```

```
## Analysis of Variance Table
##
## Response: price
##           Df      Sum Sq    Mean Sq F value    Pr(>F)
## size       1 1.5036e+13 1.5036e+13  468.56 < 2.2e-16 ***
## Residuals 298 9.5626e+12 3.2089e+10
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Use car package when type III sums of squares are needed.

```
library(car)
mod2 <- lm(price ~ size + location, data = home_prices)
anova(mod2) # default R method
```

```
## Analysis of Variance Table
##
## Response: price
##           Df      Sum Sq    Mean Sq F value    Pr(>F)
## size       1 1.5036e+13 1.5036e+13 1020.909 < 2.2e-16 ***
## location   26 5.5567e+12 2.1372e+11   14.511 < 2.2e-16 ***
## Residuals 272 4.0059e+12 1.4728e+10
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Anova(mod2, type = "III") # type II SOS, same as SPSS
```

```
## Anova Table (Type III tests)
##
## Response: price
##              Sum Sq  Df F value    Pr(>F)
## (Intercept) 1.3598e+11   1    9.233 0.002608 **
## size        9.5428e+12   1 647.954 < 2.2e-16 ***
## location    5.5567e+12  26  14.511 < 2.2e-16 ***
## Residuals   4.0059e+12 272
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Diagnostic plots.

Confidence intervals for the Mean of Y, at every X.

8 Frequencies, Proportions, Risks

PropCIs package does most of what we want

chisquare

9 Precision and Planning

MBESS has functions for this. Functions start with ss prefix. eg ss.aipe.smd

10 To do

10.1 Cat's Eye Picture

10.2 MOE for means ?